

# FIT3181: Deep Learning (2024) - Assignment 2 (Transformers)

*CE/Lecturer (Clayton):* **Dr Trung Le** | [email protected]

*Lecturer (Clayton):* **Prof Dinh Phung** | [email protected]

*Lecturer (Malaysia):* **Dr Arghya Pal** | [email protected]

*Lecturer (Malaysia):* **Dr Lim Chern Hong** | [email protected]

*Head Tutor 3181:* **Miss Vy Vo** | [[email protected] ]

*Head Tutor 5215:* **Dr Van Nguyen** | [[email protected] ]

Faculty of Information Technology, Monash University, Australia \*\*\*

## Student Information

Surname: **Ali Khan**

Firstname: **Mohib**

Student ID: **33370311**

Email: [email protected]

Your tutorial time: **Friday 10PM**

\*\*\*

## Assignment 2 – Deep Learning for Sequential Data

**Due: 11:55pm Sunday, 27 October 2024 (FIT3181)**

**Important note:** This is an individual assignment. It contributes 15% to your final mark. Read the assignment instructions carefully.

## Assignment 2's Organization

This assignment 2 has two (2) sections:

- Section 1: Fundamentals of RNNs (10 marks).
- Section 2: Deep Learning for Sequential Data (90 marks). This section is further divided into 4 parts.

The assignment 2 is organized in three (3) notebooks.

- Notebook 1 ([link](#)) [Total: 30 marks] includes Section 1 as well as Part 1 and Part 2 of Section 2.
- Notebook 2 ([link](#)) [Total: 40 marks] includes Part 3 of Section 2.
- Notebook 3 (this notebook) [Total: 30 marks] includes Part 4 of Section 2.

## What to submit

This assignment is to be completed individually and submitted to Moodle unit site. **By the due date, you are required to submit one **single zip file, named xxx\_assignment02\_solution.zip** where xxx is your student ID, to the corresponding Assignment (Dropbox) in Moodle.** You can use Google Colab to do Assignment 2 but you need to save it to an \*.ipynb file to submit to the unit Moodle.

**More importantly, if you use Google Colab to do this assignment, you need to first make a copy of this notebook on your Google drive.**

***For example, if your student ID is **12356**, then gather all of your assignment solutions to a folder, create a zip file named **123456\_assignment02\_solution.zip** and submit this file.***

Within this zip folder, you **must** submit the following files for each part:

1. **FIT3181\_DeepLearning\_Assignment2\_Official[Main].ipynb**: this is your Python notebook solution source file.
2. **FIT3181\_DeepLearning\_Assignment2\_Official[Main].html**: this is the output of your Python notebook solution *exported* in HTML format.
3. **FIT3181\_DeepLearning\_Assignment2\_Official[RNNs].ipynb**
4. **FIT3181\_DeepLearning\_Assignment2\_Official[RNNs].html**
5. **FIT3181\_DeepLearning\_Assignment2\_Official[Transformers].ipynb**
6. **FIT3181\_DeepLearning\_Assignment2\_Official[Transformers].html**
7. Any **extra files or folder** needed to complete your assignment (e.g., images used in your answers).

## Section 2: Deep Learning for Sequential Data

### Set random seeds

We need to install the package datasets for creating BERT datasets.

```
!pip install datasets
```

```
import os
import torch
import random
import requests
import pandas as pd
import numpy as np
import torch.nn as nn
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
from transformers import BertTokenizer
import os
from six.moves.urllib.request import urlretrieve
from sklearn import preprocessing
import matplotlib.pyplot as plt
plt.style.use('ggplot')

def seed_all(seed=1029):
    random.seed(seed)
    os.environ['PYTHONHASHSEED'] = str(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.cuda.manual_seed_all(seed) # if you are using multi-GPU.
    torch.backends.cudnn.benchmark = False
    torch.backends.cudnn.deterministic = True
seed_all(seed=1234)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

def seed_all(seed=1029):
    random.seed(seed)
```

```

os.environ['PYTHONHASHSEED'] = str(seed)
np.random.seed(seed)
torch.manual_seed(seed)
torch.cuda.manual_seed(seed)
torch.cuda.manual_seed_all(seed) # if you are using multi-GPU.
torch.backends.cudnn.benchmark = False
torch.backends.cudnn.deterministic = True
seed_all(seed=1234)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

```

The dataset we use **for** this assignment **is** a question classification dataset **for** which the training **set** consists of \$5,500\$ questions belonging to 6 coarse question categories including:

- abbreviation (ABBR),
- entity (ENTY),
- description (DESC),
- human (HUM),
- location (LOC) **and**
- numeric (NUM).

In this assignment, we will utilize a subset of this dataset, containing \$2,000\$ questions **for** training **and** validation. We will use 80% of those 2000 questions **for** training **and** the rest **for** validation.

**class** DataManager:

"""

*This class manages and preprocesses a simple text dataset for a sentence classification task.*

*Attributes:*

*verbose (bool): Controls verbosity for printing information during data processing.*

*max\_sentence\_len (int): The maximum length of a sentence in the dataset.*

*str\_questions (list): A list to store the string representations of the questions in the dataset.*

*str\_labels (list): A list to store the string representations of the labels in the dataset.*

*numeral\_labels (list): A list to store the numerical representations of the labels in the dataset.*

*maxlen (int): Maximum length for padding sequences. Sequences longer than this length will be truncated, and sequences shorter than this length will be padded with zeros. Defaults to 50.*

*numeral\_data (list): A list to store the numerical representations of the questions in the dataset.*

*random\_state (int): Seed value for random number generation to ensure reproducibility.*

*Set this value to a specific integer to reproduce the same random sequence every time. Defaults to 6789.*

*random (np.random.RandomState): Random number generator object initialized with the given random\_state.*

*It is used for various random operations in the class.*

#### Methods:

*maybe\_download(dir\_name, file\_name, url, verbose=True):*

*Downloads a file from a given URL if it does not exist in the specified directory.*

*The directory and file are created if they do not exist.*

*read\_data(dir\_name, file\_names):*

*Reads data from files in a directory, preprocesses it, and computes the maximum sentence length.*

*Each file is expected to contain rows in the format "<label>:<question>".*

*The labels and questions are stored as string representations.*

*manipulate\_data():*

*Performs data manipulation by tokenizing, numericalizing, and padding the text data.*

*The questions are tokenized and converted into numerical sequences using a tokenizer.*

*The sequences are padded or truncated to the maximum sequence length.*

*train\_valid\_test\_split(train\_ratio=0.9):*

*Splits the data into training, validation, and test sets based on a given ratio.*

*The data is randomly shuffled, and the specified ratio is used to determine the size of the training set.*

*The string questions, numerical data, and numerical labels are split accordingly.*

*TensorFlow `Dataset` objects are created for the training and validation sets.*

====

```
def __init__(self, verbose=True, random_state=6789):
    self.verbose = verbose
    self.max_sentence_len = 0
    self.str_questions = list()
    self.str_labels = list()
```

```

self.numeral_labels = list()
self.numeral_data = list()
self.random_state = random_state
self.random = np.random.RandomState(random_state)

@staticmethod
def maybe_download(dir_name, file_name, url, verbose=True):
    if not os.path.exists(dir_name):
        os.mkdir(dir_name)
    if not os.path.exists(os.path.join(dir_name, file_name)):
        urlretrieve(url + file_name, os.path.join(dir_name,
file_name))
    if verbose:
        print("Downloaded successfully {}".format(file_name))

def read_data(self, dir_name, file_names):
    self.str_questions = list()
    self.str_labels = list()
    for file_name in file_names:
        file_path= os.path.join(dir_name, file_name)
        with open(file_path, "r", encoding="latin-1") as f:
            for row in f:
                row_str = row.split(":")
                label, question = row_str[0], row_str[1]
                question = question.lower()
                self.str_labels.append(label)
                self.str_questions.append(question[0:-1])
                if self.max_sentence_len <
len(self.str_questions[-1]):
                    self.max_sentence_len =
len(self.str_questions[-1])

# turns labels into numbers
le = preprocessing.LabelEncoder()
le.fit(self.str_labels)
self.numeral_labels = np.array(le.transform(self.str_labels))
self.str_classes = le.classes_
self.num_classes = len(self.str_classes)
if self.verbose:
    print("\nSample questions and corresponding labels... \n")
    print(self.str_questions[0:5])
    print(self.str_labels[0:5])

def manipulate_data(self):

```

```

self.tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
vocab = self.tokenizer.get_vocab()
self.word2idx = {w: i for i, w in enumerate(vocab)}
self.idx2word = {i:w for w,i in self.word2idx.items()}
self.vocab_size = len(self.word2idx)

token_ids = []
num_seqs = []
for text in self.str_questions: # iterate over the list of text
    text_seqs = self.tokenizer.tokenize(str(text)) # tokenize each text individually
    # Convert tokens to IDs
    token_ids = self.tokenizer.convert_tokens_to_ids(text_seqs)
    # Convert token IDs to a tensor of indices using your word2idx mapping
    seq_tensor = torch.LongTensor(token_ids)
    num_seqs.append(seq_tensor) # append the tensor for each sequence

# Pad the sequences and create a tensor
if num_seqs:
    self.numeral_data = pad_sequence(num_seqs, batch_first=True)
    # Pads to max length of the sequences
    self.num_sentences, self.max_seq_len = self.numeral_data.shape

def train_valid_test_split(self, train_ratio=0.8, test_ratio = 0.1):
    train_size = int(self.num_sentences*train_ratio) +1
    test_size = int(self.num_sentences*test_ratio) +1
    valid_size = self.num_sentences - (train_size + test_size)
    data_indices = list(range(self.num_sentences))
    random.shuffle(data_indices)
    self.train_str_questions = [self.str_questions[i] for i in data_indices[:train_size]]
    self.train_numeral_labels = self.numeral_labels[data_indices[:train_size]]
    train_set_data = self.numeral_data[data_indices[:train_size]]
    train_set_labels = self.numeral_labels[data_indices[:train_size]]
    train_set_labels = torch.from_numpy(train_set_labels)
    train_set = torch.utils.data.TensorDataset(train_set_data, train_set_labels)
    self.test_str_questions = [self.str_questions[i] for i in data_indices[-test_size:]]
    self.test_numeral_labels = self.numeral_labels[data_indices[-test_size:]]
    test_set_data = self.numeral_data[data_indices[-test_size:]]

```

```

test_set_labels = self.numeral_labels[data_indices[-
test_size:]]
test_set_labels = torch.from_numpy(test_set_labels)
test_set = torch.utils.data.TensorDataset(test_set_data,
test_set_labels)
self.valid_str_questions = [self.str_questions[i] for i in
data_indices[train_size:-test_size]]
self.valid_numeral_labels =
self.numeral_labels[data_indices[train_size:-test_size]]
valid_set_data = self.numeral_data[data_indices[train_size:-
test_size]]
valid_set_labels =
self.numeral_labels[data_indices[train_size:-test_size]]
valid_set_labels = torch.from_numpy(valid_set_labels)
valid_set = torch.utils.data.TensorDataset(valid_set_data,
valid_set_labels)
self.train_loader = DataLoader(train_set, batch_size=64,
shuffle=True)
self.test_loader = DataLoader(test_set, batch_size=64,
shuffle=False)
self.valid_loader = DataLoader(valid_set, batch_size=64,
shuffle=False)

```

```

print('Loading data...')
DataManager.maybe_download("data", "train_2000.label",
"http://cogcomp.org/Data/QA/QC/")

```

```

dm = DataManager()
dm.read_data("data/", ["train_2000.label"])

```

```

print('Loading data...')
DataManager.maybe_download("data", "train_2000.label",
"http://cogcomp.org/Data/QA/QC/")

```

```

dm = DataManager()
dm.read_data("data/", ["train_2000.label"])

```

Loading data...

Downloaded successfully train\_2000.label

Sample questions and corresponding labels...



```
['manner how did serfdom develop in and then leave russia ?', 'cremat
  what films featured the character popeye doyle ?', "manner how
  can i find a list of celebrities ' real names ?", 'animal what
  fowl grabs the spotlight after the chinese year of the monkey
  ?', 'exp what is the full form of .com ?']
['DESC', 'ENTY', 'DESC', 'ENTY', 'ABBR']
```

```
dm.manipulate_data()
dm.train_valid_test_split(train_ratio=0.8, test_ratio = 0.1)
```

```
for x, y in dm.train_loader:
    print(x.shape, y.shape)
    break
```

We now declare the `BaseTrainer` **class**, which will be used later to train the subsequent deep learning models **for** text data.

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
class BaseTrainer:
    def __init__(self, model, criterion, optimizer, train_loader,
               val_loader):
        self.model = model
        self.criterion = criterion #the loss function
        self.optimizer = optimizer #the optimizer
        self.train_loader = train_loader #the train loader
        self.val_loader = val_loader #the valid loader

    #the function to train the model in many epochs
    def fit(self, num_epochs):
        self.num_batches = len(self.train_loader)

        for epoch in range(num_epochs):
            print(f'Epoch {epoch + 1}/{num_epochs}')
            train_loss, train_accuracy = self.train_one_epoch()
            val_loss, val_accuracy = self.validate_one_epoch()
            print(
                f'{self.num_batches}/{self.num_batches} - train_loss:
                {train_loss:.4f} - train_accuracy: {train_accuracy*100:.4f}% \
                - val_loss: {val_loss:.4f} - val_accuracy:
                {val_accuracy*100:.4f}%'
            )

    #train in one epoch, return the train_acc, train_loss
```

```
def train_one_epoch(self):
    self.model.train()
    running_loss, correct, total = 0.0, 0, 0
    for i, data in enumerate(self.train_loader):
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        self.optimizer.zero_grad()
        outputs = self.model(inputs)
        loss = self.criterion(outputs, labels)
        loss.backward()
        self.optimizer.step()

        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    train_accuracy = correct / total
    train_loss = running_loss / self.num_batches
    return train_loss, train_accuracy

#evaluate on a loader and return the loss and accuracy
def evaluate(self, loader):
    self.model.eval()
    loss, correct, total = 0.0, 0, 0
    with torch.no_grad():
        for data in loader:
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = self.model(inputs)
            loss = self.criterion(outputs, labels)
            loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    accuracy = correct / total
    loss = loss / len(self.val_loader)
    return loss, accuracy

#return the val_acc, val_loss, be called at the end of each epoch
def validate_one_epoch(self):
    val_loss, val_accuracy = self.evaluate(self.val_loader)
    return val_loss, val_accuracy
```

Cell In[49], line 50

The dataset we use for this assignment is a question classification dataset for which the training set consists of \$5,500\$ questions belonging to 6 coarse question categories including:

^

SyntaxError: invalid syntax

We start with importing PyTorch and NumPy and setting random seeds for PyTorch and NumPy. You can use any seeds you prefer.

```
import os
import torch
import random
import requests
import pandas as pd
import numpy as np
import torch.nn as nn
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
from transformers import BertTokenizer
import os
from six.moves.urllib.request import urlretrieve
from sklearn import preprocessing
import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

```
/Users/mohibalikhan/Desktop/Deep Learning
A02/myenv/lib/python3.9/site-packages/urllib3/__init__.py:35:
NotOpenSSLWarning: urllib3 v2 only supports OpenSSL 1.1.1+, currently
the 'ssl' module is compiled with 'LibreSSL 2.8.3'. See:
https://github.com/urllib3/urllib3/issues/3020
```

```
warnings.warn(
```

```
/Users/mohibalikhan/Desktop/Deep Learning
A02/myenv/lib/python3.9/site-packages/tqdm/auto.py:21: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
```

```
from .autonotebook import tqdm as notebook_tqdm
```

```
def seed_all(seed=1029):
    random.seed(seed)
    os.environ['PYTHONHASHSEED'] = str(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
```

```

torch.cuda.manual_seed_all(seed) # if you are using multi-GPU.
torch.backends.cudnn.benchmark = False
torch.backends.cudnn.deterministic = True
seed_all(seed=1234)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

```

## Download and preprocess the data

The dataset we use for this assignment is a question classification dataset for which the training set consists of 5,500 questions belonging to 6 coarse question categories including:

- abbreviation (ABBR),
- entity (ENTY),
- description (DESC),
- human (HUM),
- location (LOC) and
- numeric (NUM).

In this assignment, we will utilize a subset of this dataset, containing 2,000 questions for training and validation. We will use 80% of those 2000 questions for training and the rest for validation.

Preprocessing data is a crucial initial step in any machine learning or deep learning project. The *TextDataManager* class simplifies the process by providing functionalities to download and preprocess data specifically designed for the subsequent questions in this assignment. It is highly recommended to gain a comprehensive understanding of the class's functionality by **carefully reading** the content provided in the *TextDataManager* class before proceeding to answer the questions.

```
class DataManager:
```

```
    """
```

```
    This class manages and preprocesses a simple text dataset for a
    sentence classification task.
```

```
    Attributes:
```

```
        verbose (bool): Controls verbosity for printing information
        during data processing.
```

```
        max_sentence_len (int): The maximum length of a sentence in
        the dataset.
```

```
        str_questions (list): A list to store the string
        representations of the questions in the dataset.
```

```
        str_labels (list): A list to store the string representations
        of the labels in the dataset.
```

*numeral\_labels (list): A list to store the numerical representations of the labels in the dataset.*

*maxlen (int): Maximum length for padding sequences. Sequences longer than this length will be truncated, and sequences shorter than this length will be padded with zeros. Defaults to 50.*

*numeral\_data (list): A list to store the numerical representations of the questions in the dataset.*

*random\_state (int): Seed value for random number generation to ensure reproducibility.*

*Set this value to a specific integer to reproduce the same random sequence every time. Defaults to 6789.*

*random (np.random.RandomState): Random number generator object initialized with the given random\_state.*

*It is used for various random operations in the class.*

#### Methods:

*maybe\_download(dir\_name, file\_name, url, verbose=True):*

*Downloads a file from a given URL if it does not exist in the specified directory.*

*The directory and file are created if they do not exist.*

*read\_data(dir\_name, file\_names):*

*Reads data from files in a directory, preprocesses it, and computes the maximum sentence length.*

*Each file is expected to contain rows in the format "<label>:<question>".*

*The labels and questions are stored as string representations.*

*manipulate\_data():*

*Performs data manipulation by tokenizing, numericalizing, and padding the text data.*

*The questions are tokenized and converted into numerical sequences using a tokenizer.*

*The sequences are padded or truncated to the maximum sequence length.*

*train\_valid\_test\_split(train\_ratio=0.9):*

*Splits the data into training, validation, and test sets based on a given ratio.*

*The data is randomly shuffled, and the specified ratio is used to determine the size of the training set.*

*The string questions, numerical data, and numerical labels are split accordingly.*

*TensorFlow `Dataset` objects are created for the training and validation sets.*

||||

```

def __init__(self, verbose=True, random_state=6789):
    self.verbose = verbose
    self.max_sentence_len = 0
    self.str_questions = list()
    self.str_labels = list()
    self.numeral_labels = list()
    self.numeral_data = list()
    self.random_state = random_state
    self.random = np.random.RandomState(random_state)

    @staticmethod
    def maybe_download(dir_name, file_name, url, verbose=True):
        if not os.path.exists(dir_name):
            os.mkdir(dir_name)
        if not os.path.exists(os.path.join(dir_name, file_name)):
            urlretrieve(url + file_name, os.path.join(dir_name,
            file_name))
        if verbose:
            print("Downloaded successfully {}".format(file_name))

    def read_data(self, dir_name, file_names):
        self.str_questions = list()
        self.str_labels = list()
        for file_name in file_names:
            file_path= os.path.join(dir_name, file_name)
            with open(file_path, "r", encoding="latin-1") as f:
                for row in f:
                    row_str = row.split(":")
                    label, question = row_str[0], row_str[1]
                    question = question.lower()
                    self.str_labels.append(label)
                    self.str_questions.append(question[0:-1])
                    if self.max_sentence_len <
len(self.str_questions[-1]):
                        self.max_sentence_len =
len(self.str_questions[-1])

        # turns labels into numbers
        le = preprocessing.LabelEncoder()
        le.fit(self.str_labels)
        self.numeral_labels = np.array(le.transform(self.str_labels))
        self.str_classes = le.classes_
        self.num_classes = len(self.str_classes)
        if self.verbose:
            print("\nSample questions and corresponding labels... \n")

```

```

        print(self.str_questions[0:5])
        print(self.str_labels[0:5])

def manipulate_data(self):
    self.tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
    vocab = self.tokenizer.get_vocab()
    self.word2idx = {w: i for i, w in enumerate(vocab)}
    self.idx2word = {i:w for w,i in self.word2idx.items()}
    self.vocab_size = len(self.word2idx)

    token_ids = []
    num_seqs = []
    for text in self.str_questions: # iterate over the list of text
        text_seqs = self.tokenizer.tokenize(str(text)) # tokenize each text individually
        # Convert tokens to IDs
        token_ids = self.tokenizer.convert_tokens_to_ids(text_seqs)
        # Convert token IDs to a tensor of indices using your word2idx mapping
        seq_tensor = torch.LongTensor(token_ids)
        num_seqs.append(seq_tensor) # append the tensor for each sequence

    # Pad the sequences and create a tensor
    if num_seqs:
        self.numeral_data = pad_sequence(num_seqs, batch_first=True)
        # Pads to max length of the sequences
        self.num_sentences, self.max_seq_len = self.numeral_data.shape

def train_valid_test_split(self, train_ratio=0.8, test_ratio = 0.1):
    train_size = int(self.num_sentences*train_ratio) +1
    test_size = int(self.num_sentences*test_ratio) +1
    valid_size = self.num_sentences - (train_size + test_size)
    data_indices = list(range(self.num_sentences))
    random.shuffle(data_indices)
    self.train_str_questions = [self.str_questions[i] for i in data_indices[:train_size]]
    self.train_numeral_labels = self.numeral_labels[data_indices[:train_size]]
    train_set_data = self.numeral_data[data_indices[:train_size]]
    train_set_labels = self.numeral_labels[data_indices[:train_size]]
    train_set_labels = torch.from_numpy(train_set_labels)
    train_set = torch.utils.data.TensorDataset(train_set_data, train_set_labels)

```

```

self.test_str_questions = [self.str_questions[i] for i in
data_indices[-test_size:]]
self.test_numeral_labels = self.numeral_labels[data_indices[-
test_size:]]

test_set_data = self.numeral_data[data_indices[-test_size:]]
test_set_labels = self.numeral_labels[data_indices[-
test_size:]]

test_set_labels = torch.from_numpy(test_set_labels)
test_set = torch.utils.data.TensorDataset(test_set_data,
test_set_labels)
self.valid_str_questions = [self.str_questions[i] for i in
data_indices[train_size:-test_size]]
self.valid_numeral_labels =
self.numeral_labels[data_indices[train_size:-test_size]]
valid_set_data = self.numeral_data[data_indices[train_size:-
test_size]]
valid_set_labels =
self.numeral_labels[data_indices[train_size:-test_size]]
valid_set_labels = torch.from_numpy(valid_set_labels)
valid_set = torch.utils.data.TensorDataset(valid_set_data,
valid_set_labels)
self.train_loader = DataLoader(train_set, batch_size=64,
shuffle=True)
self.test_loader = DataLoader(test_set, batch_size=64,
shuffle=False)
self.valid_loader = DataLoader(valid_set, batch_size=64,
shuffle=False)

print('Loading data...')
DataManager.maybe_download("data", "train_2000.label",
"http://cogcomp.org/Data/QA/QC/")

dm = DataManager()
dm.read_data("data/", ["train_2000.label"])

Loading data...
Downloaded successfully train_2000.label

Sample questions and corresponding labels...

['manner how did serfdom develop in and then leave russia ?', 'cremat
what films featured the character popeye doyle ?', "manner how can i
find a list of celebrities ' real names ?", 'animal what fowl grabs
the spotlight after the chinese year of the monkey ?', 'exp what is
the full form of .com ?']
['DESC', 'ENTY', 'DESC', 'ENTY', 'ABBR']

dm.manipulate_data()
dm.train_valid_test_split(train_ratio=0.8, test_ratio = 0.1)

```



```

for x, y in dm.train_loader:
    print(x.shape, y.shape)
    break

```

```
torch.Size([64, 36]) torch.Size([64])
```

We now declare the BaseTrainer class, which will be used later to train the subsequent deep learning models for text data.

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```

class BaseTrainer:
    def __init__(self, model, criterion, optimizer, train_loader,
                 val_loader):
        self.model = model
        self.criterion = criterion #the loss function
        self.optimizer = optimizer #the optimizer
        self.train_loader = train_loader #the train loader
        self.val_loader = val_loader #the valid loader

#the function to train the model in many epochs
    def fit(self, num_epochs):
        self.num_batches = len(self.train_loader)

        for epoch in range(num_epochs):
            print(f'Epoch {epoch + 1}/{num_epochs}')
            train_loss, train_accuracy = self.train_one_epoch()
            val_loss, val_accuracy = self.validate_one_epoch()
            print(
                f'{self.num_batches}/{self.num_batches} - train_loss:
                {train_loss:.4f} - train_accuracy: {train_accuracy*100:.4f}% \
                - val_loss: {val_loss:.4f} - val_accuracy:
                {val_accuracy*100:.4f}%'
            )

#train in one epoch, return the train_acc, train_loss
    def train_one_epoch(self):
        self.model.train()
        running_loss, correct, total = 0.0, 0, 0
        for i, data in enumerate(self.train_loader):
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)
            self.optimizer.zero_grad()
            outputs = self.model(inputs)
            loss = self.criterion(outputs, labels)
            loss.backward()
            self.optimizer.step()

```

```
        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    train_accuracy = correct / total
    train_loss = running_loss / self.num_batches
    return train_loss, train_accuracy

#evaluate on a loader and return the loss and accuracy
def evaluate(self, loader):
    self.model.eval()
    loss, correct, total = 0.0, 0, 0
    with torch.no_grad():
        for data in loader:
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = self.model(inputs)
            loss = self.criterion(outputs, labels)
            loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    accuracy = correct / total
    loss = loss / len(self.val_loader)
    return loss, accuracy

#return the val_acc, val_loss, be called at the end of each epoch
def validate_one_epoch(self):
    val_loss, val_accuracy = self.evaluate(self.val_loader)
    return val_loss, val_accuracy
```

## Part 4: Transformer-based models for sequence modeling and neural embedding

[Total marks for this part: 30 marks]

## Question 4.1

**Implement the multi-head attention module of the Transformer for the text classification problem. The provided code is from our tutorial. In this part, we only use the output of the Transformer encoder for the classification task. For further information on the Transformer model, refer to [this paper](#).**

[Total marks for this part: 10 marks]

Below is the code of MultiHeadSelfAttention, PositionWiseFeedForward, PositionalEncoding, and EncoderLayer.

```
class MultiHeadAttention(nn.Module):
    def __init__(self, d_model, num_heads):
        super(MultiHeadAttention, self).__init__()
        # Ensure that the model dimension (d_model) is divisible by
        # the number of heads
        assert d_model % num_heads == 0, "d_model must be divisible by
        num_heads"

        # Initialize dimensions
        self.d_model = d_model # Model's dimension
        self.num_heads = num_heads # Number of attention heads
        self.d_k = d_model // num_heads # Dimension of each head's
        key, query, and value

        # Linear layers for transforming inputs
        self.W_q = nn.Linear(d_model, d_model) # Query transformation
        self.W_k = nn.Linear(d_model, d_model) # Key transformation
        self.W_v = nn.Linear(d_model, d_model) # Value transformation
        self.W_o = nn.Linear(d_model, d_model) # Output transformation

    def scaled_dot_product_attention(self, Q, K, V):
        # Calculate attention scores
        attn_scores = torch.matmul(Q, K.transpose(-2, -1)) /
        math.sqrt(self.d_k)

        # Apply mask if provided (useful for preventing attention to
        # certain parts like padding)
        #if mask is not None:
            #attn_scores = attn_scores.masked_fill(mask == 0, -1e9)

        # Softmax is applied to obtain attention probabilities
        attn_probs = torch.softmax(attn_scores, dim=-1)

        # Multiply by values to obtain the final output
```

```

        output = torch.matmul(attn_probs, V)
        return output

    def split_heads(self, x):
        # Reshape the input to have num_heads for multi-head attention
        batch_size, seq_length, d_model = x.size()
        return x.view(batch_size, seq_length, self.num_heads,
            self.d_k).transpose(1, 2)

    def combine_heads(self, x):
        # Combine the multiple heads back to original shape
        batch_size, _, seq_length, d_k = x.size()
        return x.transpose(1, 2).contiguous().view(batch_size,
            seq_length, self.d_model)

    def forward(self, Q, K, V):
        # Apply linear transformations and split heads
        Q = self.split_heads(self.W_q(Q))
        K = self.split_heads(self.W_k(K))
        V = self.split_heads(self.W_v(V))

        # Perform scaled dot-product attention
        attn_output = self.scaled_dot_product_attention(Q, K, V)

        # Combine heads and apply output transformation
        output = self.W_o(self.combine_heads(attn_output))
        return output

class PositionWiseFeedForward(nn.Module):
    def __init__(self, d_model, d_ff):
        super(PositionWiseFeedForward, self).__init__()
        self.fc1 = nn.Linear(d_model, d_ff)
        self.fc2 = nn.Linear(d_ff, d_model)
        self.relu = nn.ReLU()

    def forward(self, x):
        return self.fc2(self.relu(self.fc1(x)))

import math

class PositionalEncoding(nn.Module):
    def __init__(self, d_model, max_seq_length):
        super(PositionalEncoding, self).__init__()

        pe = torch.zeros(max_seq_length, d_model)

```

```

position = torch.arange(0, max_seq_length,
dtype=torch.float).unsqueeze(1)
div_term = torch.exp(torch.arange(0, d_model, 2).float() * -
(math.log(10000.0) / d_model))

pe[:, 0::2] = torch.sin(position * div_term)
pe[:, 1::2] = torch.cos(position * div_term)

self.register_buffer('pe', pe.unsqueeze(0))

def forward(self, x):
    return x + self.pe[:, :x.size(1)]

class EncoderLayer(nn.Module):
    def __init__(self, d_model, num_heads, d_ff, dropout):
        super(EncoderLayer, self).__init__()
        self.self_attn = MultiHeadAttention(d_model, num_heads)
        self.feed_forward = PositionWiseFeedForward(d_model, d_ff)
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x):
        attn_output = self.self_attn(x, x, x)
        x = self.norm1(x + self.dropout(attn_output))
        ff_output = self.feed_forward(x)
        x = self.norm2(x + self.dropout(ff_output))
        return x

```

Your task is to develop TransformerClassifier in which we input the embedding with the shape [batch\_size, seq\_len, embed\_dim] to some EncoderLayer (i.e., num\_layers specifies the number of EncoderLayer) and then compute the average of all token embeddings (i.e., [batch\_size, seq\_len, embed\_dim]) across the seq\_len. Finally, on the top of this average embedding, we build up a linear layer for making predictions.

```

class TransformerClassifier(nn.Module):
    def __init__(self, embed_dim, num_heads, ff_dim, num_layers,
dropout_rate=0.2, data_manager = None):
        super(TransformerClassifier, self).__init__()
        self.vocab_size = data_manager.vocab_size
        self.num_classes = data_manager.num_classes
        self.embed_dim = embed_dim
        self.max_seq_len = data_manager.max_seq_len
        self.num_heads = num_heads
        self.ff_dim = ff_dim

```

```

self.num_layers = num_layers
self.dropout_rate = dropout_rate

def build(self):
    #Insert your code here

def forward(self, x):
    #Insert your code here

class TransformerClassifier(nn.Module):
    def __init__(self, embed_dim, num_heads, ff_dim, num_layers,
                  dropout_rate=0.2, data_manager=None):
        """
        Initializes the TransformerClassifier.

        Args:
            embed_dim (int): Dimension of the input embeddings
                           (d_model).
            num_heads (int): Number of attention heads.
            ff_dim (int): Dimension of the feed-forward network.
            num_layers (int): Number of encoder layers.
            dropout_rate (float): Dropout probability.
            data_manager (DataManager): Instance containing dataset
                                     information.
        """
        super(TransformerClassifier, self).__init__()
        self.vocab_size = data_manager.vocab_size
        self.num_classes = data_manager.num_classes
        self.embed_dim = embed_dim
        self.max_seq_len = data_manager.max_seq_len
        self.num_heads = num_heads
        self.ff_dim = ff_dim
        self.num_layers = num_layers
        self.dropout_rate = dropout_rate

    def build(self):
        """
        Builds the TransformerClassifier architecture by initializing:
        - Embedding layer
        - Positional encoding
        - Encoder layers

```

```

- Final classification layer
"""

# Initialize Embedding Layer
self.embedding = nn.Embedding(self.vocab_size, self.embed_dim)

# Initialize Positional Encoding with d_model = embed_dim
self.positional_encoding =
    PositionalEncoding(d_model=self.embed_dim,
                        max_seq_length=self.max_seq_len)

# Initialize a stack of Encoder Layers
self.encoder_layers = nn.ModuleList([
    EncoderLayer(d_model=self.embed_dim,
                 num_heads=self.num_heads, d_ff=self.ff_dim,
                 dropout=self.dropout_rate)
    for _ in range(self.num_layers)
])

# Initialize the final linear layer for classification
self.fc = nn.Linear(self.embed_dim, self.num_classes)

def forward(self, x):
    """
    Defines the forward pass of the TransformerClassifier.

    Args:
        x (torch.Tensor): Input tensor of shape [batch_size,
        seq_len]

    Returns:
        torch.Tensor: Logits tensor of shape [batch_size,
        num_classes]
    """
    # Apply Embedding
    x = self.embedding(x) # Shape: [batch_size, seq_len,
    embed_dim]

    # Apply Positional Encoding
    x = self.positional_encoding(x) # Shape: [batch_size,
    seq_len, embed_dim]

    # Pass through each Encoder Layer
    for encoder in self.encoder_layers:
        x = encoder(x) # Shape remains [batch_size, seq_len,
        embed_dim]

```

```

# Compute the average of all token embeddings across the
sequence length
x = torch.mean(x, dim=1) # Shape: [batch_size, embed_dim]

# Pass the averaged embedding through the final linear layer
logits = self.fc(x) # Shape: [batch_size, num_classes]

return logits

transformer = TransformerClassifier(embed_dim=512, num_heads=8,
    ff_dim=2048, num_layers=12, dropout_rate=0.1, data_manager=
    dm)
transformer.build()
transformer = transformer.to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(transformer.parameters(), lr=1e-4, betas=
    (0.9, 0.98), eps=1e-9)
trainer = BaseTrainer(model= transformer, criterion=criterion,
    optimizer=optimizer, train_loader=dm.train_loader,
    val_loader=dm.valid_loader)
trainer.fit(num_epochs=30)

Epoch 1/30
26/26 - train_loss: 1.9181 - train_accuracy: 21.2367%
- val_loss: 0.9617 - val_accuracy: 12.6263%
Epoch 2/30
26/26 - train_loss: 1.7264 - train_accuracy: 21.6115%
- val_loss: 0.7609 - val_accuracy: 27.2727%
Epoch 3/30
26/26 - train_loss: 1.6812 - train_accuracy: 23.6727%
- val_loss: 0.8164 - val_accuracy: 26.2626%
Epoch 4/30
26/26 - train_loss: 1.6847 - train_accuracy: 22.2361%
- val_loss: 0.8395 - val_accuracy: 26.2626%
Epoch 5/30
26/26 - train_loss: 1.6815 - train_accuracy: 21.4241%
- val_loss: 0.7077 - val_accuracy: 27.2727%
Epoch 6/30
26/26 - train_loss: 1.6683 - train_accuracy: 20.5497%
- val_loss: 0.7249 - val_accuracy: 27.2727%
Epoch 7/30
26/26 - train_loss: 1.6892 - train_accuracy: 22.6109%
- val_loss: 0.7372 - val_accuracy: 27.2727%
Epoch 8/30
26/26 - train_loss: 1.3032 - train_accuracy: 45.2842%
- val_loss: 0.4200 - val_accuracy: 60.6061%

```



Epoch 9/30

26/26 - train\_loss: 0.7795 - train\_accuracy: 69.6440%  
- val\_loss: 0.2182 - val\_accuracy: 85.8586%

Epoch 10/30

26/26 - train\_loss: 0.3069 - train\_accuracy: 89.2567%  
- val\_loss: 0.0751 - val\_accuracy: 92.4242%

Epoch 11/30

26/26 - train\_loss: 0.1498 - train\_accuracy: 94.5034%  
- val\_loss: 0.0174 - val\_accuracy: 94.9495%

Epoch 12/30

26/26 - train\_loss: 0.0997 - train\_accuracy: 96.0025%  
- val\_loss: 0.0052 - val\_accuracy: 94.4444%

Epoch 13/30

26/26 - train\_loss: 0.1182 - train\_accuracy: 95.7527%  
- val\_loss: 0.0180 - val\_accuracy: 95.4545%

Epoch 14/30

26/26 - train\_loss: 0.0753 - train\_accuracy: 97.3142%  
- val\_loss: 0.0036 - val\_accuracy: 94.9495%

Epoch 15/30

26/26 - train\_loss: 0.1091 - train\_accuracy: 97.5640%  
- val\_loss: 0.0271 - val\_accuracy: 94.4444%

Epoch 16/30

26/26 - train\_loss: 0.1283 - train\_accuracy: 97.7514%  
- val\_loss: 0.0131 - val\_accuracy: 93.9394%

Epoch 17/30

26/26 - train\_loss: 0.3086 - train\_accuracy: 90.1312%  
- val\_loss: 0.0166 - val\_accuracy: 93.9394%

Epoch 18/30

26/26 - train\_loss: 0.1368 - train\_accuracy: 95.0031%  
- val\_loss: 0.0212 - val\_accuracy: 93.9394%

Epoch 19/30

26/26 - train\_loss: 0.1322 - train\_accuracy: 95.0656%  
- val\_loss: 0.0264 - val\_accuracy: 93.9394%

Epoch 20/30

26/26 - train\_loss: 0.1409 - train\_accuracy: 94.8157%  
- val\_loss: 0.0936 - val\_accuracy: 93.4343%

Epoch 21/30

26/26 - train\_loss: 0.1337 - train\_accuracy: 94.8782%  
- val\_loss: 0.0107 - val\_accuracy: 93.9394%

Epoch 22/30

26/26 - train\_loss: 0.1215 - train\_accuracy: 95.3154%  
- val\_loss: 0.0034 - val\_accuracy: 95.4545%

Epoch 23/30

```
26/26 - train_loss: 0.0941 - train_accuracy: 97.5640%
- val_loss: 0.0091 - val_accuracy: 95.9596%
Epoch 24/30
26/26 - train_loss: 0.0449 - train_accuracy: 98.4385%
- val_loss: 0.0193 - val_accuracy: 94.9495%
Epoch 25/30
26/26 - train_loss: 0.0680 - train_accuracy: 97.9388%
- val_loss: 0.0022 - val_accuracy: 94.4444%
Epoch 26/30
26/26 - train_loss: 0.0487 - train_accuracy: 98.6883%
- val_loss: 0.0008 - val_accuracy: 95.9596%
Epoch 27/30
26/26 - train_loss: 0.0314 - train_accuracy: 99.0631%
- val_loss: 0.0008 - val_accuracy: 91.4141%
Epoch 28/30
26/26 - train_loss: 0.0219 - train_accuracy: 99.4379%
- val_loss: 0.0005 - val_accuracy: 95.4545%
Epoch 29/30
26/26 - train_loss: 0.0475 - train_accuracy: 98.8132%
- val_loss: 0.0008 - val_accuracy: 95.9596%
Epoch 30/30
26/26 - train_loss: 0.0367 - train_accuracy: 99.2505%
- val_loss: 0.0001 - val_accuracy: 95.9596%
```

## Question 4.2

**Prefix prompt-tuning with Transformers: You need to implement the prefix prompt-tuning with Transformers. Basically, we base on a pre-trained Transformer, add prefix prompts, and do fine-tuning for a target dataset.**

[Total marks for this part: 10 marks]

To implement prefix prompt-tuning with pretrained Transformers, we first need to create the Bert dataset.

```
from transformers import AutoModel, AutoTokenizer, AdamW
from datasets import Dataset

model_name = "bert-base-uncased" # BERT or any similar model

# Tokenize input and prepare model inputs
tokenizer = AutoTokenizer.from_pretrained(model_name)
```

```

dataset = Dataset.from_dict({"text": dm.str_questions, "label":
                             dm.numeral_labels})

# Tokenize the dataset
def tokenize_function(examples):
    return tokenizer(examples["text"], padding="max_length",
                     truncation=True, max_length= 36)

dataset = dataset.map(tokenize_function, batched=True)
dataset.set_format(type="torch", columns=["input_ids",
                                           "attention_mask", "label"])
print(dataset)

Map: 100%|██████████| 2000/2000 [00:00<00:00, 44370.55 examples/s]

Dataset({
  features: ['text', 'label', 'input_ids', 'token_type_ids',
            'attention_mask'],
  num_rows: 2000
})

```

The following function splits the BERT dataset dataset into three BERT datasets for training, valid, and testing.

```

def train_valid_test_split(dataset, train_ratio=0.8, test_ratio =
    0.1):
    num_sentences = len(dataset)
    train_size = int(num_sentences*train_ratio) +1
    test_size = int(num_sentences*test_ratio) +1
    valid_size = num_sentences - (train_size + test_size)
    train_set = dataset[:train_size]
    train_set = Dataset.from_dict(train_set)
    train_set.set_format(type="torch", columns=["input_ids",
                                                "attention_mask", "label"])
    test_set = dataset[-test_size:]
    test_set = Dataset.from_dict(test_set)
    test_set.set_format(type="torch", columns=["input_ids",
                                                "attention_mask", "label"])
    valid_set = dataset[train_size:-test_size]
    valid_set = Dataset.from_dict(valid_set)
    valid_set.set_format(type="torch", columns=["input_ids",
                                                "attention_mask", "label"])
    train_loader = DataLoader(train_set, batch_size=64, shuffle=True)
    test_loader = DataLoader(test_set, batch_size=64, shuffle=False)
    valid_loader = DataLoader(valid_set, batch_size=64, shuffle=False)
    return train_loader, test_loader, valid_loader

```

```
train_loader, test_loader, valid_loader =
    train_valid_test_split(dataset)
```

You need to implement the class `PrefixTuningForClassification` for the prefix prompt fine-tuning. We first load a pre-trained BERT model specified by `model_name`. The parameter `prefix_length` specifies the length of the prefix prompts we add to the pre-trained BERT model. Specifically, given the input batch `[batch_size, seq_len]`, we input to the embedding layer of the pre-trained BERT model to obtain `[batch_size, seq_len, embed_size]`. We create the prefix prompts  $P$  of the size `[prefix_length, embed_size]` and concatenate to the embeddings from the pre-trained BERT to obtain `[batch_size, seq_len + prefix_length, embed_size]`. This concatenation tensor will then be fed to the encoder layers of the pre-trained BERT layer to obtain the last `[batch_size, seq_len + prefix_length, embed_size]`.

We then take mean across the `seq_len` to obtain `[batch_size, embed_size]` on which we can build up a linear layer for making predictions. Please note that **the parameters to tune include the prefix prompts  $P$  and the output linear layer**, and you should freeze the parameters of the BERT pre-trained model. Moreover, your code should cover the edge case when `prefix_length=None`. In this case, we do not insert any prefix prompts and we only do fine-tuning for the output linear layer on top.

```
class PrefixTuningForClassification(nn.Module):
    def __init__(self, model_name, prefix_length=None, data_manager =
        None):
        super(PrefixTuningForClassification, self).__init__()

        # Load the pretrained transformer model (BERT-like model)
        self.model = AutoModel.from_pretrained(model_name).to(device)
        self.hidden_size = self.model.config.hidden_size
        self.prefix_length = prefix_length
        self.num_classes = data_manager.num_classes
        # Insert your code here

    def forward(self, input_ids, attention_mask):
        # Insert your code here

import torch
import torch.nn as nn
from transformers import AutoModel, AutoTokenizer
import math

class PrefixTuningForClassification(nn.Module):
    def __init__(self, model_name, prefix_length=None,
        data_manager=None):
```

```
"""
```

*Initializes the PrefixTuningForClassification model.*

*Args:*

*model\_name (str): Name of the pre-trained Transformer model (e.g., 'bert-base-uncased').*  
*prefix\_length (int, optional): Length of the prefix prompts. If None, no prefix is added.*  
*data\_manager (DataManager): Instance containing dataset information such as vocab\_size, num\_classes, and max\_seq\_len.*

```
"""
```

```
super(PrefixTuningForClassification, self).__init__()
```

```
# Load the pretrained Transformer model (e.g., BERT)
```

```
self.model = AutoModel.from_pretrained(model_name)
```

```
self.model.to(device)
```

```
# Extract hidden size from the model configuration
```

```
self.hidden_size = self.model.config.hidden_size
```

```
# Set prefix length and number of classes from the data manager
```

```
self.prefix_length = prefix_length
```

```
self.num_classes = data_manager.num_classes
```

```
# Freeze all parameters of the pre-trained model to prevent them from being updated
```

```
for param in self.model.parameters():
```

```
    param.requires_grad = False
```

```
if self.prefix_length is not None:
```

```
    # Initialize prefix prompts as learnable parameters
```

```
    # Shape: [prefix_length, hidden_size]
```

```
    self.prefix_embeddings =
```

```
nn.Parameter(torch.randn(self.prefix_length, self.hidden_size))
```

```
# Define the final classification layer
```

```
# This layer maps the averaged embeddings to the number of classes
```

```
self.classifier = nn.Linear(self.hidden_size, self.num_classes)
```

```
def build(self):
```

```
"""
```

*Prepares the model for training by moving it to the appropriate device.*

*This method can be expanded if additional setup is required.*

```

"""
# Currently, all setup is done in __init__
pass

def forward(self, input_ids, attention_mask):
    """
    Defines the forward pass of the model.

    Args:
        input_ids (torch.Tensor): Tensor of shape [batch_size,
        seq_len] containing token IDs.
        attention_mask (torch.Tensor): Tensor of shape
        [batch_size, seq_len] containing attention masks.

    Returns:
        torch.Tensor: Logits tensor of shape [batch_size,
        num_classes].
    """

    batch_size, seq_len = input_ids.size()

    if self.prefix_length is not None:
        # Expand prefix prompts to match the batch size
        # Shape after expansion: [batch_size, prefix_length,
        hidden_size]
        prefix =
self.prefix_embeddings.unsqueeze(0).expand(batch_size, -1, -1)

        # Obtain input embeddings from the pre-trained model's
        embedding layer
        # Shape: [batch_size, seq_len, hidden_size]
        inputs_embeds = self.model.embeddings(input_ids)

        # Concatenate prefix prompts with input embeddings
        # Shape: [batch_size, prefix_length + seq_len,
        hidden_size]
        concatenated_embeds = torch.cat((prefix, inputs_embeds),
        dim=1)

        # Create a new attention mask that accounts for the prefix
        prompts
        # Prefix tokens have an attention mask of 1
        # Original attention_mask shape: [batch_size, seq_len]
        # New attention_mask shape: [batch_size, prefix_length +
        seq_len]
        prefix_attention_mask = torch.ones(batch_size,
        self.prefix_length).to(device)
        concatenated_attention_mask =
        torch.cat((prefix_attention_mask, attention_mask), dim=1)

```

```

        # Pass the concatenated embeddings and the new attention
        mask to the Transformer encoder
        outputs = self.model(
            inputs_embeds=concatenated_embeds,
            attention_mask=concatenated_attention_mask
        )

        # Extract the last hidden state
        # Shape: [batch_size, prefix_length + seq_len,
        hidden_size]
        last_hidden_state = outputs.last_hidden_state
    else:
        # If no prefix prompts, proceed as standard Transformer
        classification
        outputs = self.model(
            input_ids=input_ids,
            attention_mask=attention_mask
        )
        last_hidden_state = outputs.last_hidden_state

    # Compute the mean of the token embeddings across the sequence
    length
    # Shape: [batch_size, hidden_size]
    # This includes prefix prompts if they were added
    pooled_output = torch.mean(last_hidden_state, dim=1)

    # Pass the pooled output through the classification layer to
    obtain logits
    # Shape: [batch_size, num_classes]
    logits = self.classifier(pooled_output)

    return logits

```

You can use the following FineTunedBaseTrainer to train the prompt fine-tuning models.

```

class FineTunedBaseTrainer:
    def __init__(self, model, criterion, optimizer, train_loader,
        val_loader):
        self.model = model
        self.criterion = criterion #the loss function
        self.optimizer = optimizer #the optimizer
        self.train_loader = train_loader #the train loader
        self.val_loader = val_loader #the valid loader

    #the function to train the model in many epochs
    def fit(self, num_epochs):

```

```

self.num_batches = len(self.train_loader)

for epoch in range(num_epochs):
    print(f'Epoch {epoch + 1}/{num_epochs}')
    train_loss, train_accuracy = self.train_one_epoch()
    val_loss, val_accuracy = self.validate_one_epoch()
    print(
        f'{self.num_batches}/{self.num_batches} - train_loss:
        {train_loss:.4f} - train_accuracy: {train_accuracy*100:.4f}% \
        - val_loss: {val_loss:.4f} - val_accuracy:
        {val_accuracy*100:.4f}%'
    )

#train in one epoch, return the train_acc, train_loss
def train_one_epoch(self):
    self.model.train()
    running_loss, correct, total = 0.0, 0, 0
    for batch in self.train_loader:
        input_ids = batch["input_ids"].to(device)
        attention_mask = batch["attention_mask"].to(device)
        labels = batch["label"].to(device)
        self.optimizer.zero_grad()
        outputs = self.model(input_ids= input_ids, attention_mask=
attention_mask)
        loss = self.criterion(outputs, labels)
        loss.backward()
        self.optimizer.step()

        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    train_accuracy = correct / total
    train_loss = running_loss / self.num_batches
    return train_loss, train_accuracy

#evaluate on a loader and return the loss and accuracy
def evaluate(self, loader):
    self.model.eval()
    loss, correct, total = 0.0, 0, 0
    with torch.no_grad():
        for batch in loader:
            input_ids = batch["input_ids"].to(device)
            labels = batch["label"].to(device)
            attention_mask = batch["attention_mask"].to(device)

```



```

        outputs = self.model(input_ids= input_ids,
                              attention_mask= attention_mask)
        loss = self.criterion(outputs, labels)
        loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    accuracy = correct / total
    loss = loss / len(self.val_loader)
    return loss, accuracy

```

*#return the val\_acc, val\_loss, be called at the end of each epoch*

```

def validate_one_epoch(self):
    val_loss, val_accuracy = self.evaluate(self.val_loader)
    return val_loss, val_accuracy

```

We declare and train the prefix-prompt tuning model. In addition, you need to be patient with this model because it might converge slowly with many epochs.

```

prefix_tuning_model = PrefixTuningForClassification(model_name =
    "bert-base-uncased", prefix_length = 5, data_manager =
    dm).to(device)

if prefix_tuning_model.prefix_length is not None:
    optimizer =
        torch.optim.Adam(list(prefix_tuning_model.classifier.parameters())
        + [prefix_tuning_model.prefix_embeddings], lr=5e-5)
else:
    optimizer =
        torch.optim.Adam(prefix_tuning_model.classifier.parameters(),
        lr=1e-4)

criterion = nn.CrossEntropyLoss()
trainer = FineTunedBaseTrainer(model= prefix_tuning_model,
    criterion=criterion, optimizer=optimizer,
    train_loader=train_loader, val_loader=valid_loader)
trainer.fit(num_epochs=100)

```

Epoch 1/100

26/26 – train\_loss: 1.7966 – train\_accuracy: 16.5522%  
 – val\_loss: 0.8773 – val\_accuracy: 16.6667%

Epoch 2/100

26/26 – train\_loss: 1.7542 – train\_accuracy: 19.2380%  
 – val\_loss: 0.8589 – val\_accuracy: 21.7172%

Epoch 3/100

26/26 – train\_loss: 1.7190 – train\_accuracy: 23.8601%  
 – val\_loss: 0.8476 – val\_accuracy: 26.2626%

Epoch 4/100

26/26 - train\_loss: 1.6917 - train\_accuracy: 28.2324%  
- val\_loss: 0.8372 - val\_accuracy: 32.3232%  
Epoch 5/100  
26/26 - train\_loss: 1.6776 - train\_accuracy: 32.4172%  
- val\_loss: 0.8295 - val\_accuracy: 36.8687%  
Epoch 6/100  
26/26 - train\_loss: 1.6574 - train\_accuracy: 35.0406%  
- val\_loss: 0.8192 - val\_accuracy: 41.4141%  
Epoch 7/100  
26/26 - train\_loss: 1.6335 - train\_accuracy: 38.4760%  
- val\_loss: 0.8094 - val\_accuracy: 44.4444%  
Epoch 8/100  
26/26 - train\_loss: 1.6178 - train\_accuracy: 39.2255%  
- val\_loss: 0.7982 - val\_accuracy: 46.4646%  
Epoch 9/100  
26/26 - train\_loss: 1.5995 - train\_accuracy: 40.7870%  
- val\_loss: 0.7920 - val\_accuracy: 46.4646%  
Epoch 10/100  
26/26 - train\_loss: 1.5868 - train\_accuracy: 41.7864%  
- val\_loss: 0.7822 - val\_accuracy: 46.9697%  
Epoch 11/100  
26/26 - train\_loss: 1.5754 - train\_accuracy: 42.9731%  
- val\_loss: 0.7730 - val\_accuracy: 46.9697%  
Epoch 12/100  
26/26 - train\_loss: 1.5610 - train\_accuracy: 43.4728%  
- val\_loss: 0.7715 - val\_accuracy: 47.4747%  
Epoch 13/100  
26/26 - train\_loss: 1.5509 - train\_accuracy: 47.0956%  
- val\_loss: 0.7635 - val\_accuracy: 47.9798%  
Epoch 14/100  
26/26 - train\_loss: 1.5447 - train\_accuracy: 47.6577%  
- val\_loss: 0.7624 - val\_accuracy: 49.4949%  
Epoch 15/100  
26/26 - train\_loss: 1.5377 - train\_accuracy: 49.7189%  
- val\_loss: 0.7588 - val\_accuracy: 52.0202%  
Epoch 16/100  
26/26 - train\_loss: 1.5217 - train\_accuracy: 49.9063%  
- val\_loss: 0.7525 - val\_accuracy: 53.0303%  
Epoch 17/100  
26/26 - train\_loss: 1.5820 - train\_accuracy: 52.4047%  
- val\_loss: 0.7518 - val\_accuracy: 53.5354%  
Epoch 18/100  
26/26 - train\_loss: 1.4978 - train\_accuracy: 51.9051%

– val\_loss: 0.7473 – val\_accuracy: 54.0404%  
Epoch 19/100  
26/26 – train\_loss: 1.4993 – train\_accuracy: 51.4678%  
– val\_loss: 0.7416 – val\_accuracy: 54.0404%  
Epoch 20/100  
26/26 – train\_loss: 1.4805 – train\_accuracy: 51.9051%  
– val\_loss: 0.7319 – val\_accuracy: 53.5354%  
Epoch 21/100  
26/26 – train\_loss: 1.4799 – train\_accuracy: 52.2798%  
– val\_loss: 0.7276 – val\_accuracy: 56.5657%  
Epoch 22/100  
26/26 – train\_loss: 1.4791 – train\_accuracy: 53.9663%  
– val\_loss: 0.7217 – val\_accuracy: 58.5859%  
Epoch 23/100  
26/26 – train\_loss: 1.4569 – train\_accuracy: 57.8389%  
– val\_loss: 0.7166 – val\_accuracy: 60.1010%  
Epoch 24/100  
26/26 – train\_loss: 1.4457 – train\_accuracy: 57.4641%  
– val\_loss: 0.7129 – val\_accuracy: 61.1111%  
Epoch 25/100  
26/26 – train\_loss: 1.4454 – train\_accuracy: 58.1512%  
– val\_loss: 0.7129 – val\_accuracy: 61.1111%  
Epoch 26/100  
26/26 – train\_loss: 1.4318 – train\_accuracy: 58.3385%  
– val\_loss: 0.7069 – val\_accuracy: 64.6465%  
Epoch 27/100  
26/26 – train\_loss: 1.4994 – train\_accuracy: 61.0868%  
– val\_loss: 0.6965 – val\_accuracy: 65.1515%  
Epoch 28/100  
26/26 – train\_loss: 1.4105 – train\_accuracy: 61.2742%  
– val\_loss: 0.6933 – val\_accuracy: 65.1515%  
Epoch 29/100  
26/26 – train\_loss: 1.4078 – train\_accuracy: 61.3367%  
– val\_loss: 0.6913 – val\_accuracy: 64.1414%  
Epoch 30/100  
26/26 – train\_loss: 1.3973 – train\_accuracy: 59.9001%  
– val\_loss: 0.6882 – val\_accuracy: 64.6465%  
Epoch 31/100  
26/26 – train\_loss: 1.3943 – train\_accuracy: 61.0244%  
– val\_loss: 0.6848 – val\_accuracy: 64.6465%  
Epoch 32/100  
26/26 – train\_loss: 1.3743 – train\_accuracy: 61.2117%  
– val\_loss: 0.6749 – val\_accuracy: 65.1515%

Epoch 33/100

26/26 - train\_loss: 1.3664 - train\_accuracy: 60.9619%  
- val\_loss: 0.6666 - val\_accuracy: 66.1616%

Epoch 34/100

26/26 - train\_loss: 1.3710 - train\_accuracy: 62.8982%  
- val\_loss: 0.6599 - val\_accuracy: 67.6768%

Epoch 35/100

26/26 - train\_loss: 1.3650 - train\_accuracy: 63.7726%  
- val\_loss: 0.6567 - val\_accuracy: 69.6970%

Epoch 36/100

26/26 - train\_loss: 1.3419 - train\_accuracy: 63.8351%  
- val\_loss: 0.6553 - val\_accuracy: 69.6970%

Epoch 37/100

26/26 - train\_loss: 1.3314 - train\_accuracy: 64.4597%  
- val\_loss: 0.6481 - val\_accuracy: 71.2121%

Epoch 38/100

26/26 - train\_loss: 1.3270 - train\_accuracy: 66.7083%  
- val\_loss: 0.6426 - val\_accuracy: 71.7172%

Epoch 39/100

26/26 - train\_loss: 1.3310 - train\_accuracy: 65.0843%  
- val\_loss: 0.6360 - val\_accuracy: 72.7273%

Epoch 40/100

26/26 - train\_loss: 1.3201 - train\_accuracy: 65.7714%  
- val\_loss: 0.6290 - val\_accuracy: 72.7273%

Epoch 41/100

26/26 - train\_loss: 1.3227 - train\_accuracy: 66.3960%  
- val\_loss: 0.6320 - val\_accuracy: 72.2222%

Epoch 42/100

26/26 - train\_loss: 1.3270 - train\_accuracy: 66.8332%  
- val\_loss: 0.6280 - val\_accuracy: 71.7172%

Epoch 43/100

26/26 - train\_loss: 1.2870 - train\_accuracy: 69.0194%  
- val\_loss: 0.6221 - val\_accuracy: 73.7374%

Epoch 44/100

26/26 - train\_loss: 1.3164 - train\_accuracy: 68.6446%  
- val\_loss: 0.6156 - val\_accuracy: 73.7374%

Epoch 45/100

26/26 - train\_loss: 1.3032 - train\_accuracy: 69.3941%  
- val\_loss: 0.6155 - val\_accuracy: 75.2525%

Epoch 46/100

26/26 - train\_loss: 1.2816 - train\_accuracy: 69.0194%  
- val\_loss: 0.6114 - val\_accuracy: 76.7677%

Epoch 47/100

26/26 - train\_loss: 1.2757 - train\_accuracy: 69.8938%  
- val\_loss: 0.6127 - val\_accuracy: 75.2525%  
Epoch 48/100  
26/26 - train\_loss: 1.2547 - train\_accuracy: 69.8938%  
- val\_loss: 0.6116 - val\_accuracy: 76.7677%  
Epoch 49/100  
26/26 - train\_loss: 1.3526 - train\_accuracy: 71.0806%  
- val\_loss: 0.6058 - val\_accuracy: 76.2626%  
Epoch 50/100  
26/26 - train\_loss: 1.2515 - train\_accuracy: 70.7683%  
- val\_loss: 0.6004 - val\_accuracy: 76.7677%  
Epoch 51/100  
26/26 - train\_loss: 1.2482 - train\_accuracy: 71.8926%  
- val\_loss: 0.5958 - val\_accuracy: 77.7778%  
Epoch 52/100  
26/26 - train\_loss: 1.2546 - train\_accuracy: 71.3929%  
- val\_loss: 0.5920 - val\_accuracy: 77.7778%  
Epoch 53/100  
26/26 - train\_loss: 1.2356 - train\_accuracy: 72.1424%  
- val\_loss: 0.5877 - val\_accuracy: 77.7778%  
Epoch 54/100  
26/26 - train\_loss: 1.2351 - train\_accuracy: 71.0181%  
- val\_loss: 0.5809 - val\_accuracy: 77.2727%  
Epoch 55/100  
26/26 - train\_loss: 1.2142 - train\_accuracy: 72.5172%  
- val\_loss: 0.5795 - val\_accuracy: 77.7778%  
Epoch 56/100  
26/26 - train\_loss: 1.2214 - train\_accuracy: 71.7676%  
- val\_loss: 0.5728 - val\_accuracy: 77.7778%  
Epoch 57/100  
26/26 - train\_loss: 1.1934 - train\_accuracy: 72.5796%  
- val\_loss: 0.5736 - val\_accuracy: 78.2828%  
Epoch 58/100  
26/26 - train\_loss: 1.1982 - train\_accuracy: 73.2042%  
- val\_loss: 0.5741 - val\_accuracy: 78.2828%  
Epoch 59/100  
26/26 - train\_loss: 1.2049 - train\_accuracy: 73.2042%  
- val\_loss: 0.5685 - val\_accuracy: 78.2828%  
Epoch 60/100  
26/26 - train\_loss: 1.1991 - train\_accuracy: 72.2049%  
- val\_loss: 0.5663 - val\_accuracy: 79.7980%  
Epoch 61/100  
26/26 - train\_loss: 1.1827 - train\_accuracy: 73.8913%

- val\_loss: 0.5608 - val\_accuracy: 80.3030%  
Epoch 62/100  
26/26 - train\_loss: 1.1840 - train\_accuracy: 73.7664%  
- val\_loss: 0.5586 - val\_accuracy: 79.7980%  
Epoch 63/100  
26/26 - train\_loss: 1.1793 - train\_accuracy: 72.7046%  
- val\_loss: 0.5544 - val\_accuracy: 80.3030%  
Epoch 64/100  
26/26 - train\_loss: 1.2102 - train\_accuracy: 72.8919%  
- val\_loss: 0.5482 - val\_accuracy: 79.7980%  
Epoch 65/100  
26/26 - train\_loss: 1.1643 - train\_accuracy: 73.3916%  
- val\_loss: 0.5424 - val\_accuracy: 78.7879%  
Epoch 66/100  
26/26 - train\_loss: 1.1646 - train\_accuracy: 72.7670%  
- val\_loss: 0.5403 - val\_accuracy: 80.8081%  
Epoch 67/100  
26/26 - train\_loss: 1.1268 - train\_accuracy: 74.8907%  
- val\_loss: 0.5411 - val\_accuracy: 80.8081%  
Epoch 68/100  
26/26 - train\_loss: 1.1512 - train\_accuracy: 73.4541%  
- val\_loss: 0.5410 - val\_accuracy: 79.7980%  
Epoch 69/100  
26/26 - train\_loss: 1.1259 - train\_accuracy: 75.2030%  
- val\_loss: 0.5366 - val\_accuracy: 79.7980%  
Epoch 70/100  
26/26 - train\_loss: 1.1457 - train\_accuracy: 73.9538%  
- val\_loss: 0.5298 - val\_accuracy: 79.2929%  
Epoch 71/100  
26/26 - train\_loss: 1.1577 - train\_accuracy: 75.0156%  
- val\_loss: 0.5249 - val\_accuracy: 78.7879%  
Epoch 72/100  
26/26 - train\_loss: 1.1124 - train\_accuracy: 75.4528%  
- val\_loss: 0.5176 - val\_accuracy: 78.7879%  
Epoch 73/100  
26/26 - train\_loss: 1.1133 - train\_accuracy: 75.2655%  
- val\_loss: 0.5182 - val\_accuracy: 78.2828%  
Epoch 74/100  
26/26 - train\_loss: 1.1247 - train\_accuracy: 75.8901%  
- val\_loss: 0.5142 - val\_accuracy: 78.2828%  
Epoch 75/100  
26/26 - train\_loss: 1.1076 - train\_accuracy: 75.5778%  
- val\_loss: 0.5137 - val\_accuracy: 79.7980%

Epoch 76/100

26/26 - train\_loss: 1.0773 - train\_accuracy: 75.5778%  
- val\_loss: 0.5131 - val\_accuracy: 78.7879%

Epoch 77/100

26/26 - train\_loss: 1.0921 - train\_accuracy: 75.8901%  
- val\_loss: 0.5077 - val\_accuracy: 79.2929%

Epoch 78/100

26/26 - train\_loss: 1.1060 - train\_accuracy: 75.1405%  
- val\_loss: 0.5003 - val\_accuracy: 79.2929%

Epoch 79/100

26/26 - train\_loss: 1.1032 - train\_accuracy: 73.8289%  
- val\_loss: 0.4982 - val\_accuracy: 80.3030%

Epoch 80/100

26/26 - train\_loss: 1.0889 - train\_accuracy: 75.0156%  
- val\_loss: 0.4920 - val\_accuracy: 79.7980%

Epoch 81/100

26/26 - train\_loss: 1.0743 - train\_accuracy: 74.9532%  
- val\_loss: 0.4931 - val\_accuracy: 78.7879%

Epoch 82/100

26/26 - train\_loss: 1.0769 - train\_accuracy: 75.1405%  
- val\_loss: 0.4912 - val\_accuracy: 80.3030%

Epoch 83/100

26/26 - train\_loss: 1.0762 - train\_accuracy: 75.2655%  
- val\_loss: 0.4887 - val\_accuracy: 80.3030%

Epoch 84/100

26/26 - train\_loss: 1.0528 - train\_accuracy: 75.6402%  
- val\_loss: 0.4851 - val\_accuracy: 80.8081%

Epoch 85/100

26/26 - train\_loss: 1.0674 - train\_accuracy: 76.9519%  
- val\_loss: 0.4784 - val\_accuracy: 81.3131%

Epoch 86/100

26/26 - train\_loss: 1.0488 - train\_accuracy: 75.6402%  
- val\_loss: 0.4732 - val\_accuracy: 80.8081%

Epoch 87/100

26/26 - train\_loss: 1.0311 - train\_accuracy: 77.0144%  
- val\_loss: 0.4713 - val\_accuracy: 80.8081%

Epoch 88/100

26/26 - train\_loss: 1.0458 - train\_accuracy: 77.9513%  
- val\_loss: 0.4693 - val\_accuracy: 81.8182%

Epoch 89/100

26/26 - train\_loss: 1.0253 - train\_accuracy: 75.9525%  
- val\_loss: 0.4699 - val\_accuracy: 80.3030%

Epoch 90/100

```
26/26 - train_loss: 1.0454 - train_accuracy: 75.5778%
- val_loss: 0.4658 - val_accuracy: 81.3131%
Epoch 91/100
26/26 - train_loss: 1.0266 - train_accuracy: 76.3898%
- val_loss: 0.4629 - val_accuracy: 81.3131%
Epoch 92/100
26/26 - train_loss: 1.0014 - train_accuracy: 76.6396%
- val_loss: 0.4593 - val_accuracy: 81.3131%
Epoch 93/100
26/26 - train_loss: 1.0273 - train_accuracy: 76.4522%
- val_loss: 0.4563 - val_accuracy: 81.3131%
Epoch 94/100
26/26 - train_loss: 1.0082 - train_accuracy: 75.9525%
- val_loss: 0.4563 - val_accuracy: 81.8182%
Epoch 95/100
26/26 - train_loss: 1.0159 - train_accuracy: 76.0775%
- val_loss: 0.4502 - val_accuracy: 81.8182%
Epoch 96/100
26/26 - train_loss: 1.0058 - train_accuracy: 77.7639%
- val_loss: 0.4460 - val_accuracy: 81.3131%
Epoch 97/100
26/26 - train_loss: 1.0012 - train_accuracy: 77.5765%
- val_loss: 0.4430 - val_accuracy: 80.8081%
Epoch 98/100
26/26 - train_loss: 0.9897 - train_accuracy: 76.7645%
- val_loss: 0.4434 - val_accuracy: 80.8081%
Epoch 99/100
26/26 - train_loss: 0.9977 - train_accuracy: 77.2017%
- val_loss: 0.4381 - val_accuracy: 81.3131%
Epoch 100/100
26/26 - train_loss: 1.0084 - train_accuracy: 77.7639%
- val_loss: 0.4393 - val_accuracy: 80.8081%
```

```
from sklearn.metrics import accuracy_score, classification_report
```

```
# Set the model to evaluation mode
prefix_tuning_model.eval()
```

```
all_preds = []
all_labels = []
```

```
with torch.no_grad():
    for batch in test_loader:
        # Move input data to the appropriate device
```



```

input_ids = batch['input_ids'].to(device)
attention_mask = batch['attention_mask'].to(device)
labels = batch['label'].to(device)

# Forward pass to get outputs/logits
outputs = prefix_tuning_model(input_ids, attention_mask)

# Get the predicted class by taking the argmax
_, preds = torch.max(outputs, dim=1)

# Append predictions and true labels to the lists
all_preds.extend(preds.cpu().numpy())
all_labels.extend(labels.cpu().numpy())

# Compute the overall accuracy
test_accuracy = accuracy_score(all_labels, all_preds)

print(f"The Test Accuracy we obtain is: {test_accuracy * 100}%")

The Test Accuracy we obtain is: 79.1044776119403%

```

### Question 4.3

For any models defined in the previous questions (of all parts), you are free to fine-tune hyperparameters, e.g., optimizer, learning\_rate, state\_sizes, such that you get a best model, i.e., the one with the highest accuracy on the test set. You will need to report (i) what is your best model, (ii) its accuracy on the test set, and (iii) the values of its hyperparameters. Note that you must report your best model's accuracy with rounding to 4 decimal places, i.e., 0.xxxx. You will also need to upload your best model (or provide us with the link to download your best model). The assessment will be based on your best model's accuracy, with up to 9 marks available, specifically:

- The best accuracy  $\geq 0.97$ : 10 marks
- $0.97 > \text{The best accuracy} \geq 0.92$ : 7 marks
- $0.92 > \text{The best accuracy} \geq 0.85$ : 4 marks
- The best accuracy  $< 0.85$ : 0 mark

For this question, you can put below the code to train the best model. In this case, you need to show your code and the evidence of running regarding the best model. Moreover, if you save the best model, you need to provide the link to download the best model, the code to load the best model, and then evaluate on the test set.

## Give your answer here.

- (i) What is your best model?
- (ii) The accuracy of your best model on the test set
- (iii) The values of the hyperparameters of your best model
- (iv) The link to download your best model

### Answer

- (i) --> I have used the RNN Model from 3.2.1 as my best model
- (ii) --> The accuracy it gets is 0.98
- (iii) --> There was no need to do any modifications on my model but rather I just went with these paramters

```
cell_type='gru',  
state_sizes=[64, 128],  
output_type='max',  
data_manager= dm,  
run_mode='init-fine-tune'
```

- (iv) --> Link to download

([https://drive.google.com/file/d/19o8YpaD1CZGtTcnarIF69K2iX7ry-9Xn/view?usp=drive\\_link](https://drive.google.com/file/d/19o8YpaD1CZGtTcnarIF69K2iX7ry-9Xn/view?usp=drive_link))

```
import os  
import torch  
import random  
import requests  
import pandas as pd  
import numpy as np  
import torch.nn as nn  
from torch.utils.data import DataLoader  
from torch.nn.utils.rnn import pad_sequence  
from transformers import BertTokenizer  
import os  
from six.moves.urllib.request import urlretrieve # type: ignore  
from sklearn import preprocessing  
import matplotlib.pyplot as plt
```

```
plt.style.use('ggplot')
```

```
def seed_all(seed=1029):
    random.seed(seed)
    os.environ['PYTHONHASHSEED'] = str(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.cuda.manual_seed_all(seed) # if you are using multi-GPU.
    torch.backends.cudnn.benchmark = False
    torch.backends.cudnn.deterministic = True
seed_all(seed=1234)
```

```
class DataManager:
    """
    This class manages and preprocesses a simple text dataset for a
    sentence classification task.

    Attributes:
        verbose (bool): Controls verbosity for printing information
            during data processing.
        max_sentence_len (int): The maximum length of a sentence in
            the dataset.
        str_questions (list): A list to store the string
            representations of the questions in the dataset.
        str_labels (list): A list to store the string representations
            of the labels in the dataset.
        numeral_labels (list): A list to store the numerical
            representations of the labels in the dataset.
        numeral_data (list): A list to store the numerical
            representations of the questions in the dataset.
        random_state (int): Seed value for random number generation to
            ensure reproducibility.
            Set this value to a specific integer to reproduce the same
            random sequence every time. Defaults to 6789.
        random (np.random.RandomState): Random number generator object
            initialized with the given random_state.
            It is used for various random operations in the class.

    Methods:
        maybe_download(dir_name, file_name, url, verbose=True):
            Downloads a file from a given URL if it does not exist in
            the specified directory.
            The directory and file are created if they do not exist.

        read_data(dir_name, file_names):
            Reads data from files in a directory, preprocesses it, and
            computes the maximum sentence length.
            Each file is expected to contain rows in the format "
            <label>:<question>".
```

*The labels and questions are stored as string representations.*

*manipulate\_data():*

*Performs data manipulation by tokenizing, numericalizing, and padding the text data.*

*The questions are tokenized and converted into numerical sequences using a tokenizer.*

*The sequences are padded or truncated to the maximum sequence length.*

*train\_valid\_test\_split(train\_ratio=0.9):*

*Splits the data into training, validation, and test sets based on a given ratio.*

*The data is randomly shuffled, and the specified ratio is used to determine the size of the training set.*

*The string questions, numerical data, and numerical labels are split accordingly.*

*TensorFlow `Dataset` objects are created for the training and validation sets.*

"""

```
def __init__(self, verbose=True, random_state=6789):
    self.verbose = verbose
    self.max_sentence_len = 0
    self.str_questions = list()
    self.str_labels = list()
    self.numeral_labels = list()
    self.maxlen = None
    self.numeral_data = list()
    self.random_state = random_state
    self.random = np.random.RandomState(random_state)

    @staticmethod
    def maybe_download(dir_name, file_name, url, verbose=True):
        if not os.path.exists(dir_name):
            os.mkdir(dir_name)
        if not os.path.exists(os.path.join(dir_name, file_name)):
            urlretrieve(url + file_name, os.path.join(dir_name,
            file_name))
        if verbose:
            print("Downloaded successfully {}".format(file_name))

    def read_data(self, dir_name, file_names):
        self.str_questions = list()
```

```

self.str_labels = list()
for file_name in file_names:
    file_path= os.path.join(dir_name, file_name)
    with open(file_path, "r", encoding="latin-1") as f:
        for row in f:
            row_str = row.split(":")
            label, question = row_str[0], row_str[1]
            question = question.lower()
            self.str_labels.append(label)
            self.str_questions.append(question[0:-1])
            if self.max_sentence_len <
len(self.str_questions[-1]):
                self.max_sentence_len =
len(self.str_questions[-1])

# turns labels into numbers
le = preprocessing.LabelEncoder()
le.fit(self.str_labels)
self.numeral_labels = np.array(le.transform(self.str_labels))
self.str_classes = le.classes_
self.num_classes = len(self.str_classes)
if self.verbose:
    print("\nSample questions and corresponding labels... \n")
    print(self.str_questions[0:5])
    print(self.str_labels[0:5])

def manipulate_data(self):
    self.tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
    vocab = self.tokenizer.get_vocab()
    self.word2idx = {w: i for i, w in enumerate(vocab)}
    self.idx2word = {i:w for w,i in self.word2idx.items()}
    self.vocab_size = len(self.word2idx)

    token_ids = []
    num_seqs = []
    for text in self.str_questions: # iterate over the list of
text
        text_seqs = self.tokenizer.tokenize(str(text)) # tokenize
each text individually
        # Convert tokens to IDs
        token_ids = self.tokenizer.convert_tokens_to_ids(text_seqs)
        # Convert token IDs to a tensor of indices using your
word2idx mapping
        seq_tensor = torch.LongTensor(token_ids)

```

```

        num_seqs.append(seq_tensor) # append the tensor for each
sequence

# Pad the sequences and create a tensor
if num_seqs:
    self.numeral_data = pad_sequence(num_seqs, batch_first=True)
    # Pads to max length of the sequences
    self.num_sentences, self.maxlen = self.numeral_data.shape

def train_valid_test_split(self, train_ratio=0.8, test_ratio =
0.1):
    train_size = int(self.num_sentences*train_ratio) +1
    test_size = int(self.num_sentences*test_ratio) +1
    valid_size = self.num_sentences - (train_size + test_size)
    data_indices = list(range(self.num_sentences))
    random.shuffle(data_indices)
    self.train_str_questions = [self.str_questions[i] for i in
data_indices[:train_size]]
    self.train_numeral_labels =
self.numeral_labels[data_indices[:train_size]]
    train_set_data = self.numeral_data[data_indices[:train_size]]
    train_set_labels =
self.numeral_labels[data_indices[:train_size]]
    train_set_labels = torch.from_numpy(train_set_labels)
    train_set = torch.utils.data.TensorDataset(train_set_data,
train_set_labels)
    self.test_str_questions = [self.str_questions[i] for i in
data_indices[-test_size:]]
    self.test_numeral_labels = self.numeral_labels[data_indices[-
test_size:]]
    test_set_data = self.numeral_data[data_indices[-test_size:]]
    test_set_labels = self.numeral_labels[data_indices[-
test_size:]]
    test_set_labels = torch.from_numpy(test_set_labels)
    test_set = torch.utils.data.TensorDataset(test_set_data,
test_set_labels)
    self.valid_str_questions = [self.str_questions[i] for i in
data_indices[train_size:-test_size]]
    self.valid_numeral_labels =
self.numeral_labels[data_indices[train_size:-test_size]]
    valid_set_data = self.numeral_data[data_indices[train_size:-
test_size]]
    valid_set_labels =
self.numeral_labels[data_indices[train_size:-test_size]]
    valid_set_labels = torch.from_numpy(valid_set_labels)
    valid_set = torch.utils.data.TensorDataset(valid_set_data,
valid_set_labels)
    self.train_loader = DataLoader(train_set, batch_size=64,
shuffle=True) # you can change the batch size if needed

```

```

self.test_loader = DataLoader(test_set, batch_size=64,
                               shuffle=False) # you can change the batch size if needed
self.valid_loader = DataLoader(valid_set, batch_size=64,
                                shuffle=False) # you can change the batch size if needed

print('Loading data...')
DataManager.maybe_download("data", "train_2000.label",
                           "http://cogcomp.org/Data/QA/QC/")

dm = DataManager()
dm.read_data("data/", ["train_2000.label"])

dm.manipulate_data()
dm.train_valid_test_split(train_ratio=0.8, test_ratio = 0.1)

Loading data...
Downloaded successfully train_2000.label

Sample questions and corresponding labels...

['manner how did serfdom develop in and then leave russia ?', 'cremat
what films featured the character popeye doyle ?', "manner how can i
find a list of celebrities ' real names ?", 'animal what fowl grabs
the spotlight after the chinese year of the monkey ?', 'exp what is
the full form of .com ?']
['DESC', 'ENTY', 'DESC', 'ENTY', 'ABBR']

for x, y in dm.train_loader:
    print(x.shape, y.shape)
    break

torch.Size([64, 36]) torch.Size([64])

#device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
-- i dont need this line

import torch
# Check if MPS is available
if torch.backends.mps.is_available():
    device = torch.device("mps")
else:
    device = torch.device("cpu")

class BaseTrainer:
    def __init__(self, model, criterion, optimizer, train_loader,
                 val_loader):
        self.model = model

```

```

self.criterion = criterion #the loss function
self.optimizer = optimizer #the optimizer
self.train_loader = train_loader #the train loader
self.val_loader = val_loader #the valid loader

#the function to train the model in many epochs
def fit(self, num_epochs):
    self.num_batches = len(self.train_loader)

    for epoch in range(num_epochs):
        print(f'Epoch {epoch + 1}/{num_epochs}')
        train_loss, train_accuracy = self.train_one_epoch()
        val_loss, val_accuracy = self.validate_one_epoch()
        print(
            f'{self.num_batches}/{self.num_batches} - train_loss:
{train_loss:.4f} - train_accuracy: {train_accuracy*100:.4f}% \
- val_loss: {val_loss:.4f} - val_accuracy:
{val_accuracy*100:.4f}%'
)

#train in one epoch, return the train_acc, train_loss
def train_one_epoch(self):
    self.model.train()
    running_loss, correct, total = 0.0, 0, 0
    for i, data in enumerate(self.train_loader):
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        self.optimizer.zero_grad()
        outputs = self.model(inputs)
        loss = self.criterion(outputs, labels)
        loss.backward()
        self.optimizer.step()

        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    train_accuracy = correct / total
    train_loss = running_loss / self.num_batches
    return train_loss, train_accuracy

#evaluate on a loader and return the loss and accuracy
def evaluate(self, loader):
    self.model.eval()
    loss, correct, total = 0.0, 0, 0

```



```

with torch.no_grad():
    for data in loader:
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = self.model(inputs)
        loss = self.criterion(outputs, labels)
        loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    accuracy = correct / total
    loss = loss / len(self.val_loader)
    return loss, accuracy

#return the val_acc, val_loss, be called at the end of each epoch
def validate_one_epoch(self):
    val_loss, val_accuracy = self.evaluate(self.val_loader)
    return val_loss, val_accuracy

class BaseRNN(nn.Module):
    def __init__(self, cell_type='gru', embed_size=128, state_sizes=
        [128, 128], output_type="mean", data_manager=None):
        super().__init__()
        self.cell_type = cell_type                # Type of RNN
        cell: 'simple_rnn', 'gru', or 'lstm'
        self.state_sizes = state_sizes             # List of hidden
        sizes for each RNN layer
        self.embed_size = embed_size               # Dimension of
        word embeddings
        self.output_type = output_type             # Output strategy:
        'last_state', 'mean', or 'max'
        self.data_manager = data_manager           # Data manager
        containing dataset information
        self.vocab_size = self.data_manager.vocab_size # Size of the
        vocabulary

    # Static method to return the corresponding RNN layer based on
    cell_type
    @staticmethod
    def get_layer(cell_type='gru', input_size=128, state_size=128):
        if cell_type == 'gru':
            return nn.GRU(input_size=input_size,
                hidden_size=state_size, batch_first=True)
        elif cell_type == 'lstm':
            return nn.LSTM(input_size=input_size,
                hidden_size=state_size, batch_first=True)
        else: # 'simple_rnn'

```

```

        return nn.RNN(input_size=input_size,
                       hidden_size=state_size, batch_first=True)

def build(self):
    # Embedding layer to convert word indices to embeddings
    self.embed = nn.Embedding(self.vocab_size, self.embed_size)

    # ModuleList to hold multiple RNN layers
    self.rnn_layers = nn.ModuleList()
    input_size = self.embed_size # Initial input size is the
    embedding size

    # Create RNN layers based on state_sizes
    for state_size in self.state_sizes:
        rnn_layer = self.get_layer(self.cell_type,
                                   input_size=input_size, state_size=state_size)
        self.rnn_layers.append(rnn_layer)
        input_size = state_size # Output size of current layer is
        input size for next layer

    # Fully connected layer for classification
    self.fc = nn.Linear(self.state_sizes[-1],
                        self.data_manager.num_classes)

def forward(self, x):
    # x: input tensor of shape [batch_size, seq_len]
    # Pass input through embedding layer
    e = self.embed(x) # e: [batch_size, seq_len, embed_size]

    # Pass embeddings through the stacked RNN layers
    output = e # Initial input to the first RNN layer
    for rnn_layer in self.rnn_layers:
        if self.cell_type == 'lstm':
            output, (h_n, c_n) = rnn_layer(output)
        else:
            output, h_n = rnn_layer(output)

    # Obtain the representation based on output_type
    if self.output_type == "last_state":
        # For last_state, use the last hidden state h_n
        # h_n shape: [num_layers * num_directions, batch_size,
        hidden_size]
        h = h_n[-1] # Get the last layer's hidden state; shape:
        [batch_size, hidden_size]
    elif self.output_type == "mean":
        # For mean, average over the sequence length dimension

```

```

        h = torch.mean(output, dim=1) # h: [batch_size,
        hidden_size]
    elif self.output_type == "max":
        # For max, take the maximum over the sequence length
        dimension
        h, _ = torch.max(output, dim=1) # h: [batch_size,
        hidden_size]
    else:
        raise ValueError("Invalid output_type. Choose from
        'last_state', 'mean', or 'max'.")

    logits = self.fc(h) # logits: [batch_size, num_classes]
    return logits

class RNN(BaseRNN):
    def __init__(self, cell_type='gru', embed_size=128, state_sizes=
        [128, 128], output_type='mean', data_manager=None,
        run_mode='scratch', embed_model='glove-wiki-gigaword-
        100'):
        super().__init__(cell_type, embed_size, state_sizes,
        output_type, data_manager)
        self.run_mode = run_mode
        self.embed_model = embed_model
        if not os.path.exists("embeddings"):
            os.makedirs("embeddings")
        self.embed_path = "embeddings/E.npy"
        if self.run_mode != 'scratch':
            self.embed_size = int(self.embed_model.split("-")[-1])
        self.word2idx = data_manager.word2idx
        self.word2vect = None
        self.embed_matrix = np.zeros((self.vocab_size,
        self.embed_size))
        # Call self.build() after all attributes are set
        self.build()

    def build_embedding_matrix(self):
        # Check if the embedding matrix already exists
        if os.path.exists(self.embed_path):
            print(f"Loading embedding matrix from {self.embed_path}")
            self.embed_matrix = np.load(self.embed_path)
        else:
            # Download the pretrained embedding model
            print(f"Downloading embedding model
            {self.embed_model}...")
            self.word2vect = api.load(self.embed_model)
            print("Building embedding matrix...")
            # Initialize the embedding matrix

```

```

        for word, idx in self.word2idx.items():
            if word in self.word2vect:
                # If the word is in the pretrained embeddings, use
                # its vector
                self.embed_matrix[idx] = self.word2vect[word]
            else:
                # Otherwise, initialize a random vector
                self.embed_matrix[idx] = np.random.uniform(-0.25,
0.25, self.embed_size)
        # Save the embedding matrix for future use
        np.save(self.embed_path, self.embed_matrix)
        print(f"Saved embedding matrix to {self.embed_path}")

def build(self):
    # Build the embedding layer based on run_mode
    if self.run_mode == 'scratch':
        # Initialize the embedding layer from scratch
        self.embed = nn.Embedding(self.vocab_size,
self.embed_size)
    else:
        # Build the embedding matrix using the pretrained
        # embeddings
        self.build_embedding_matrix()
        # Create the embedding layer and load the pretrained
        # weights
        self.embed = nn.Embedding(self.vocab_size,
self.embed_size)

    self.embed.weight.data.copy_(torch.from_numpy(self.embed_matrix))
    if self.run_mode == 'init-only':
        # Freeze the embedding layer weights
        self.embed.weight.requires_grad = False
    elif self.run_mode == 'init-fine-tune':
        # Allow the embedding layer to be fine-tuned
        self.embed.weight.requires_grad = True
    else:
        raise ValueError("Invalid run_mode. Choose from
'scratch', 'init-only', or 'init-fine-tune'.")
    # Build the RNN layers and the fully connected layer
    self.rnn_layers = nn.ModuleList()
    input_size = self.embed_size
    for state_size in self.state_sizes:
        rnn_layer = self.get_layer(self.cell_type,
input_size=input_size, state_size=state_size)
        self.rnn_layers.append(rnn_layer)
        input_size = state_size

```

```

# Fully connected layer for classification
self.fc = nn.Linear(self.state_sizes[-1],
self.data_manager.num_classes)

# Insert your code here
print("\nTraining RNN with run_mode 'init-fine-tune'")

# Instantiate the RNN model
rnn_init_fine_tune = RNN(
    cell_type='gru',
    state_sizes=[64, 128],
    output_type='max',
    data_manager= dm,
    run_mode='init-fine-tune',
    embed_model='glove-wiki-gigaword-100'
)
rnn_init_fine_tune.to(device)

# Define the loss function
criterion_init_fine_tune = nn.CrossEntropyLoss()

# Define the optimizer
optimizer_init_fine_tune =
    torch.optim.Adam(rnn_init_fine_tune.parameters(), lr=0.001)

# Initialize the trainer
trainer_init_fine_tune = BaseTrainer(
    model=rnn_init_fine_tune,
    criterion=criterion_init_fine_tune,
    optimizer=optimizer_init_fine_tune,
    train_loader=dm.train_loader,
    val_loader=dm.valid_loader
)

# Train the model
trainer_init_fine_tune.fit(num_epochs=30)

# Evaluate on the validation set
val_loss_init_fine_tune, val_accuracy_init_fine_tune =
    trainer_init_fine_tune.validate_one_epoch()
print(f"Validation Loss: {val_loss_init_fine_tune:.4f} - Validation
    Accuracy: {val_accuracy_init_fine_tune*100:.2f}%")

Training RNN with run_mode 'init-fine-tune'
Loading embedding matrix from embeddings/E.npy

```

Epoch 1/30

26/26 - train\_loss: 1.5941 - train\_accuracy: 41.9738%  
- val\_loss: 0.5975 - val\_accuracy: 54.0404%

Epoch 2/30

26/26 - train\_loss: 1.0853 - train\_accuracy: 68.8944%  
- val\_loss: 0.3698 - val\_accuracy: 90.4040%

Epoch 3/30

26/26 - train\_loss: 0.4852 - train\_accuracy: 90.6933%  
- val\_loss: 0.1171 - val\_accuracy: 93.4343%

Epoch 4/30

26/26 - train\_loss: 0.2300 - train\_accuracy: 93.3791%  
- val\_loss: 0.0450 - val\_accuracy: 93.9394%

Epoch 5/30

26/26 - train\_loss: 0.1311 - train\_accuracy: 95.1280%  
- val\_loss: 0.0265 - val\_accuracy: 93.9394%

Epoch 6/30

26/26 - train\_loss: 0.0967 - train\_accuracy: 96.6896%  
- val\_loss: 0.0159 - val\_accuracy: 95.4545%

Epoch 7/30

26/26 - train\_loss: 0.0617 - train\_accuracy: 98.3760%  
- val\_loss: 0.0107 - val\_accuracy: 94.9495%

Epoch 8/30

26/26 - train\_loss: 0.0510 - train\_accuracy: 98.5634%  
- val\_loss: 0.0064 - val\_accuracy: 96.4646%

Epoch 9/30

26/26 - train\_loss: 0.0397 - train\_accuracy: 98.9382%  
- val\_loss: 0.0050 - val\_accuracy: 96.9697%

Epoch 10/30

26/26 - train\_loss: 0.0241 - train\_accuracy: 99.4379%  
- val\_loss: 0.0037 - val\_accuracy: 95.9596%

Epoch 11/30

26/26 - train\_loss: 0.0303 - train\_accuracy: 99.0631%  
- val\_loss: 0.0028 - val\_accuracy: 97.4747%

Epoch 12/30

26/26 - train\_loss: 0.0138 - train\_accuracy: 99.6877%  
- val\_loss: 0.0022 - val\_accuracy: 97.4747%

Epoch 13/30

26/26 - train\_loss: 0.0075 - train\_accuracy: 99.9375%  
- val\_loss: 0.0020 - val\_accuracy: 96.9697%

Epoch 14/30

26/26 - train\_loss: 0.0110 - train\_accuracy: 99.7502%  
- val\_loss: 0.0016 - val\_accuracy: 97.9798%

Epoch 15/30

26/26 - train\_loss: 0.0046 - train\_accuracy: 99.9375%  
- val\_loss: 0.0014 - val\_accuracy: 97.9798%  
Epoch 16/30  
26/26 - train\_loss: 0.0036 - train\_accuracy: 100.0000%  
- val\_loss: 0.0014 - val\_accuracy: 96.9697%  
Epoch 17/30  
26/26 - train\_loss: 0.0048 - train\_accuracy: 99.8751%  
- val\_loss: 0.0012 - val\_accuracy: 96.9697%  
Epoch 18/30  
26/26 - train\_loss: 0.0091 - train\_accuracy: 99.7502%  
- val\_loss: 0.0010 - val\_accuracy: 97.9798%  
Epoch 19/30  
26/26 - train\_loss: 0.0031 - train\_accuracy: 100.0000%  
- val\_loss: 0.0010 - val\_accuracy: 96.9697%  
Epoch 20/30  
26/26 - train\_loss: 0.0020 - train\_accuracy: 100.0000%  
- val\_loss: 0.0008 - val\_accuracy: 97.9798%  
Epoch 21/30  
26/26 - train\_loss: 0.0016 - train\_accuracy: 100.0000%  
- val\_loss: 0.0007 - val\_accuracy: 97.9798%  
Epoch 22/30  
26/26 - train\_loss: 0.0015 - train\_accuracy: 100.0000%  
- val\_loss: 0.0007 - val\_accuracy: 97.9798%  
Epoch 23/30  
26/26 - train\_loss: 0.0013 - train\_accuracy: 100.0000%  
- val\_loss: 0.0006 - val\_accuracy: 97.9798%  
Epoch 24/30  
26/26 - train\_loss: 0.0012 - train\_accuracy: 100.0000%  
- val\_loss: 0.0006 - val\_accuracy: 97.9798%  
Epoch 25/30  
26/26 - train\_loss: 0.0011 - train\_accuracy: 100.0000%  
- val\_loss: 0.0005 - val\_accuracy: 97.9798%  
Epoch 26/30  
26/26 - train\_loss: 0.0010 - train\_accuracy: 100.0000%  
- val\_loss: 0.0005 - val\_accuracy: 97.9798%  
Epoch 27/30  
26/26 - train\_loss: 0.0010 - train\_accuracy: 100.0000%  
- val\_loss: 0.0005 - val\_accuracy: 97.9798%  
Epoch 28/30  
26/26 - train\_loss: 0.0009 - train\_accuracy: 100.0000%  
- val\_loss: 0.0004 - val\_accuracy: 97.9798%  
Epoch 29/30  
26/26 - train\_loss: 0.0009 - train\_accuracy: 100.0000%

```
- val_loss: 0.0004 - val_accuracy: 97.9798%
```

```
Epoch 30/30
```

```
26/26 - train_loss: 0.0008 - train_accuracy: 100.0000%
```

```
- val_loss: 0.0004 - val_accuracy: 97.9798%
```

```
Validation Loss: 0.0004 - Validation Accuracy: 97.98%
```

```
trainer_init_fine_tune.fit(num_epochs=30)
```

```
test_loss, test_acc = trainer_init_fine_tune.evaluate(dm.test_loader)
print(f'test_loss: {test_loss:.4f} - test_accuracy:
      {test_acc*100:.4f}%')
```

```
test_loss: 0.4790 - test_accuracy: 98.0100%
```

```
torch.save(rnn_init_fine_tune.state_dict(), 'final_best_model.pth')
rnn_init_fine_tune.load_state_dict(torch.load('final_best_model.pth',
      map_location=device))
```

```
test_loss, test_acc = trainer_init_fine_tune.evaluate(dm.test_loader)
print(f'test_loss: {test_loss:.4f} - test_accuracy:
      {test_acc*100:.4f}%')
```

```
test_loss: 0.4790 - test_accuracy: 98.0100%
```

```
/var/folders/gm/2ndchprn0czbhd7f3zbnh2f80000gn/T/ipykernel_77091/245286
FutureWarning: You are using `torch.load` with `weights_only=False`
(the current default value), which uses the default pickle module
implicitly. It is possible to construct malicious pickle data which
will execute arbitrary code during unpickling (See
https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-
models for more details). In a future release, the default value for
`weights_only` will be flipped to `True`. This limits the functions
that could be executed during unpickling. Arbitrary objects will no
longer be allowed to be loaded via this mode unless they are
explicitly allowlisted by the user via
`torch.serialization.add_safe_globals`. We recommend you start setting
`weights_only=True` for any use case where you don't have full control
of the loaded file. Please open an issue on GitHub for any issues
related to this experimental feature.
```

```
rnn_init_fine_tune.load_state_dict(torch.load('final_best_model.pth',
      map_location=device))
```

**GOOD LUCK WITH YOUR ASSIGNMENT 2!**

**END OF ASSIGNMENT**



